

```
pip install ucimlrepo
```

```
Collecting ucimlrepo
  Downloading ucimlrepo-0.0.7-py3-none-any.whl.metadata (5.5 kB)
Requirement already satisfied: pandas>=1.0.0 in /usr/local/lib/python3.12/dist-packages (from ucimlrepo) (2.2.2)
Requirement already satisfied: certifi>=2020.12.5 in /usr/local/lib/python3.12/dist-packages (from ucimlrepo) (2026.1.4)
Requirement already satisfied: numpy>=1.26.0 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.0.0->ucimlrepo) (2.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.0.0->ucimlrepo) (2.3.1)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.0.0->ucimlrepo) (2020.1.1)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.0.0->ucimlrepo) (2022.7)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas>=1.0.0->ucimlrepo) (1.16.0)
Downloading ucimlrepo-0.0.7-py3-none-any.whl (8.0 kB)
Installing collected packages: ucimlrepo
Successfully installed ucimlrepo-0.0.7
```

```
from ucimlrepo import fetch_ucirepo

# fetch dataset
breast_cancer_wisconsin_diagnostic = fetch_ucirepo(id=17)

# data (as pandas dataframes)
X = breast_cancer_wisconsin_diagnostic.data.features
y = breast_cancer_wisconsin_diagnostic.data.targets

# metadata
print(breast_cancer_wisconsin_diagnostic.metadata)

# variable information
print(breast_cancer_wisconsin_diagnostic.variables)
```

```
{'uci_id': 17, 'name': 'Breast Cancer Wisconsin (Diagnostic)', 'repository_url': 'https://archive.ics.uci.edu/dataset/17/Breast+Cancer+Wisconsin+\(Diagnostic\).arff'}

      name    role      type demographic description units \
0        ID     ID Categorical      None      None  None
1  Diagnosis Target Categorical      None      None  None
2      radius1 Feature Continuous      None      None  None
3      texture1 Feature Continuous      None      None  None
4      perimeter1 Feature Continuous      None      None  None
5       area1 Feature Continuous      None      None  None
6   smoothness1 Feature Continuous      None      None  None
7   compactness1 Feature Continuous      None      None  None
8      concavity1 Feature Continuous      None      None  None
9  concave_points1 Feature Continuous      None      None  None
10     symmetry1 Feature Continuous      None      None  None
11 fractal_dimension1 Feature Continuous      None      None  None
12      radius2 Feature Continuous      None      None  None
13      texture2 Feature Continuous      None      None  None
14      perimeter2 Feature Continuous      None      None  None
15       area2 Feature Continuous      None      None  None
16   smoothness2 Feature Continuous      None      None  None
17   compactness2 Feature Continuous      None      None  None
18      concavity2 Feature Continuous      None      None  None
19  concave_points2 Feature Continuous      None      None  None
20      symmetry2 Feature Continuous      None      None  None
21 fractal_dimension2 Feature Continuous      None      None  None
22      radius3 Feature Continuous      None      None  None
23      texture3 Feature Continuous      None      None  None
24      perimeter3 Feature Continuous      None      None  None
25       area3 Feature Continuous      None      None  None
26   smoothness3 Feature Continuous      None      None  None
27   compactness3 Feature Continuous      None      None  None
28      concavity3 Feature Continuous      None      None  None
29  concave_points3 Feature Continuous      None      None  None
30      symmetry3 Feature Continuous      None      None  None
31 fractal_dimension3 Feature Continuous      None      None  None

  missing_values
0        no
1        no
2        no
3        no
4        no
5        no
6        no
7        no
8        no
9        no
10       no
11       no
12       no
13       no
14       no
15       no
16       no
17       no
```

```
18      no
19      no
20      no
```

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    confusion_matrix
)
```

```
# Convert y DataFrame to Series and encode
y = y.iloc[:, 0].map({'M': 1, 'B': 0})
```

```
print(y.value_counts())
```

```
Diagnosis
0    357
1    212
Name: count, dtype: int64
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,
    random_state=42,
    stratify=y
)
```

```
print("Training samples:", X_train.shape[0])
print("Test samples:", X_test.shape[0])
```

```
Training samples: 455
Test samples: 114
```

```
scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
log_model = LogisticRegression(max_iter=500)
log_model.fit(X_train_scaled, y_train)

# Predictions
y_train_pred_log = log_model.predict(X_train_scaled)
y_test_pred_log = log_model.predict(X_test_scaled)

# Errors
train_error_log = 1 - accuracy_score(y_train, y_train_pred_log)
test_error_log = 1 - accuracy_score(y_test, y_test_pred_log)

print("LOGISTIC REGRESSION RESULTS")
print("Train Error:", round(train_error_log, 4))
print("Test Error:", round(test_error_log, 4))
print("Accuracy:", round(accuracy_score(y_test, y_test_pred_log), 4))
print("Precision:", round(precision_score(y_test, y_test_pred_log), 4))
print("Recall:", round(recall_score(y_test, y_test_pred_log), 4))
print("F1-score:", round(f1_score(y_test, y_test_pred_log), 4))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_test_pred_log))
```

```
LOGISTIC REGRESSION RESULTS
Train Error: 0.0132
Test Error: 0.0351
Accuracy: 0.9649
Precision: 0.975
Recall: 0.9286
F1-score: 0.9512
Confusion Matrix:
[[71  1]
 [ 3 39]]
```

```
tree_model = DecisionTreeClassifier(random_state=42)
tree_model.fit(X_train, y_train)

# Predictions
y_train_pred_tree = tree_model.predict(X_train)
y_test_pred_tree = tree_model.predict(X_test)

# Errors
train_error_tree = 1 - accuracy_score(y_train, y_train_pred_tree)
test_error_tree = 1 - accuracy_score(y_test, y_test_pred_tree)

print("DECISION TREE RESULTS")
print("Train Error:", round(train_error_tree, 4))
print("Test Error:", round(test_error_tree, 4))
print("Accuracy:", round(accuracy_score(y_test, y_test_pred_tree), 4))
print("Precision:", round(precision_score(y_test, y_test_pred_tree), 4))
print("Recall:", round(recall_score(y_test, y_test_pred_tree), 4))
print("F1-score:", round(f1_score(y_test, y_test_pred_tree), 4))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_test_pred_tree))
```

```
DECISION TREE RESULTS
Train Error: 0.0
Test Error: 0.0702
Accuracy: 0.9298
Precision: 0.9048
Recall: 0.9048
F1-score: 0.9048
Confusion Matrix:
 [[68  4]
 [ 4 38]]
```

```
comparison = pd.DataFrame({
    "Model": ["Logistic Regression", "Decision Tree"],
    "Train Error": [train_error_log, train_error_tree],
    "Test Error": [test_error_log, test_error_tree]
})

comparison
```

	Model	Train Error	Test Error	Actions
0	Logistic Regression	0.013187	0.035088	
1	Decision Tree	0.000000	0.070175	

Next steps: [Generate code with comparison](#) [New interactive sheet](#)

Conclusion

The Logistic Regression model performs consistently on both training and test data, showing only a small difference in error. This indicates that the model generalizes well and is able to make reliable predictions on unseen cases.

On the other hand, the Decision Tree performs extremely well on the training data but shows noticeably worse performance on the test set. This gap suggests overfitting, where the model learns the training data too closely and struggles to generalize.

Based on these results, Logistic Regression is the better choice for deployment since it provides more stable and dependable performance.

Relevant ML Considerations

- Logistic Regression requires feature scaling, whereas Decision Trees do not.
- Decision Trees are more likely to overfit if not properly constrained.
- Care was taken to avoid data leakage by fitting the scaler only on training data